









# Alex Towell

- [lex@metafunctor.com](mailto:lex@metafunctor.com) 
- <https://metafunctor.com> 
- <https://github.com/queelius> 
- <https://twitter.com/queelius> 
- URL for this talk: <https://github.com/queelius/sluug-llm> 

# Outline of Talk

- Theoretical Background
- Go over a simple language model
  - $n$ -gram model (Jupyter Notebook)
  - Easy to understand and helps us understand some aspects of LLMs.
- Show an application of LLMs:
  - Try to make a database search API intelligent (NLP) with small LLMs.
- Open Discussion

# Good-Old-Fashioned AI (GOFAI)

- Find a way to symbolically represent the problem and then use logic or rules to solve it.
  - Programming 
  - Rule-based systems 
  - First-order logic
- LLMs are *good* at using these tools. 
  - Integrate Prolog with LLM tool-use to help with planning and reasoning?

# Reductive Reasoning

GOFAI works for a lot of problems we care about:

- Filter everything through our small working memory.
  - Inductive bias: Makes assumptions about the world.
  - Help us generalize out-of-distribution. 🧠
- Take big problems and break down into simpler problems.
- Solve simpler problems and combine.

# Limits of GOFAI

Many problems are hard to break down into simpler parts.

- Whole greater than the sum of its parts.
- Too complex to solve reductively.
  - We can't program computers to do it. 🧑
  - Identifying cats in pictures? 🐱
  - The hard problems are easy and the easy problems are hard.  
-- Steven Pinker
  - Playing with legos is hard but multivariate calculus is easy (for a computer).

# How Do Our Brains Work?

Brains programmed by evolution to survive in a complex world.

- It's a prediction engine: it learns to predict the world.
- The unconscious mind is not limited by a small "working memory"
- It can do things we don't understand how to do.
- Brain is a black box. (See: *Interpretable ML*)

# Machine Learning

💡 Let's have the computer learn from data.

- Since the real world is too complex, let's have the computer learn from data like we do.
- There are three main types of learning.
  - Supervised Learning (SL)
  - Unsupervised Learning
  - Reinforcement Learning (RL)
- *Spoiler:* LLMs use self-supervised learning (SSL) and RL (RLHF).

## Type of Learning (1): Supervised Learning

**Learning from labeled data.** We have some input and output data, and we want to learn how to map the input to the output.

- Given an (unknown) function  $f$  and a set of input-output pairs  $(x, f(x))$ , learn a function  $\hat{f}$  that approximates  $f$  on the input-output pairs.
- E.g., classification:  $f : [\text{🐱 or 🐶}] \mapsto \{\text{🐱}, \text{🐶}\}$ .
  - Use  $\hat{f}$  to predict 🐱 or 🐶 for new images.
- Easiest problem to solve in ML. But: limited by data.
- **Fine-Tuning** LLMs is supervised learning: improve it on specific labeled tasks.



## Type of Learning (2): Unsupervised Learning

**No labeled data.** Learn the underlying structure of the data.

- Clustering: Grouping similar data points. (See: *RAG*)
- Dimensionality Reduction: Learn *efficient* representations of the data.
  - Very hard and one of the most important problems in ML.
- Density Estimation: Stochastic estimate of process that generated the observed data. Say the process generates  $(x, y)$  pairs and we estimate its density  $\Pr(x, y)$ .
  - Classification (supervised):  $\Pr(y|x) = \Pr(x, y) / \Pr(x)$
- **Pre-training LLMs** is like unsupervised learning. Learn a good representation and probability distribution of the *raw* text using self-supervised learning (SSL).

## Final Type of Learning (3): Reinforcement Learning

This is an agentic approach to learning. Agent interacts with environment and learns from the rewards it receives.

- *Goal*: maximize the expected sum of rewards.
- *Spoiler*: Agentic frameworks that include LLMs as a prediction component is a very active area of research.
- Prediction + Search = Planning
  - Counterfactual reasoning
- Hypothesis: `Compression = Prediction = Intelligence`
- Big reason a lot of people are excited about Sora.
  - Has everyone seen the Sora videos?
  - "Intuitive" world simulation (embedded in the weights of a giant NN).

# Early Failures in ML

Reality is really complicated:  $(x_1, x_2, \dots, x_n)$ ,  
 $n$  extremely large, and each  $x_i$  is some complex object.

Early efforts in ML were not very successful.

- Overfitting, curse of dimensionality, lack of data/compute.
- To combat lack of data/compute, clever solutions developed.
  - Many of these methods are no longer around.

"The biggest lesson that can be read from 70 years of AI research is that general methods that leverage computation are ultimately the most effective, and by a large margin."

-- Richard Sutton in "The Bitter Lesson":

# Neural Networks

Neural Networks (NN) are one the solutions that stuck around.

- It fell out of favor for a while, but it's back.
- Universal function approximator.
  - Can learn to represent any function.
  - But: need a lot of data to do so and be difficult to train.
- NNs seem to scale to as much data and compute as we can throw at them.

# Inductive Bias

Observations may have an infinite set of hypothesis that are compatible with the data.

- **Inductive Bias:** The set of assumptions that the model makes about the data.
- **Occam's Razor:** choose the simplest hypothesis that is compatible with the data. (See *Solomonoff Induction*.)
- Generalizing out-of-distribution (OOD) from inputs not in the training data.
- **Problem:** We are almost *always* out-of-distribution.
  - Except in toy problems (see: early successes)
- Good inductive biases are necessary for generalization.
- **No Free Lunch Theorem:** No model is optimal for all tasks.

# Era of Deep Learning

One of the hardest parts is learning sample efficient representation of the data.

- Layers of NN learn progressively higher-level representations: Pixels -> Edges -> Objects
- AlexNet (2012) was the first to show that deep learning could work well on large-scale datasets.

## **Era of Deep Learning (cont.)**

DNNs (feed-forward) learn little circuit programs that can generate parts of the training data. (Image stolen from Jeff Dean's slides.)

- Hundreds of layers: can learn pretty complicated programs.
- (What a human can do in a half a second, a DNN can do?)

## **Era of Generative AI**

Generative AI "reverses" the arrows  
- Image to text, image to image,  
etc.

- They learn something about the data generating process (DGP).
- They have completely changed our expectations of what computers can do.



## Era of Generative AI (cont.)

We now have computers that can see, hear, understand, and generate all of these things.

Let's go look at **Sora**: generative video, or world(s) simulator?

- **Scaling**: And increasing the scale (data, compute) increase their capabilities. See: Scaling laws.
  - Need a lot more *compute*.
  - It's going to get wild(er).
  - Hypothesis: `Prediction = Compression = Intelligence`.

# Large Language Models (LLMs)

Autoregressive (AR) models learn a probability distribution over training data by using self-supervised learning (SSL):

$$\Pr\{x_1, x_2, \dots, x_T\} = \prod_{t=1}^T \Pr\{x_t | x_1, \dots, x_{t-1}\}$$

- This is hard to learn, but with enough data and compute, a lot seems possible.
- LLMs have a nice advantage since language is designed to have a very low dimensionality and have a high signal to noise ratio.
  - Representation learning is easier in language than in other domains.
    - Still learns representations ( word2vec )
- **Language** represents much of the things that humans care and think about, so learning to predict it is a kind of general intelligence. (See: Sparks of AGI by Microsoft)

# Sampling from LLMs

There are many different ways to sample from LLMs and change the behavior of the model.

- **Temperature:** Rescaling the logits before applying the softmax function.
  - $T = 1$ : estimates the probability distribution.
  - $T < 1$ : reduces randomness, i.e., more predictable outputs.
  - $T > 1$ : increases randomness, i.e., more unpredictable outputs.

Good for controlling *exploitation vs exploration* if repeatedly sampling from the model to generate new or different outputs.

- **Top- $k$  and Top- $p$  Sampling:** Choose the top- $k$  or top- $p$  tokens and sample from them.
- **Beam Search:** Explore multiple paths and sample the most likely.

# Prompting Strategies

Early models were very sensitive to the prompt.

- Makes sense, they were trained to generate the data.
- If you condition on crazy data, you get crazy outputs.

$$\Pr\{\text{more crazy}|\text{crazy}\}$$

Various prompting strategies have been developed to help the model generate more reliable outputs:

- Chain-of-thought (CoT)
- Tree-of-thought (ToT)
- and so on

- Pre-trained models predict the raw data distribution.
- Fine-Tuning: Train the model to a specific data that is more relevant to a task.
- RLHF: Bias the model to produce outputs that people prefer.
- **Goal:** Enable the generation of new data points that resemble in some ways the characteristics the data.
  - At inference, outputs are almost always out-of-distribution (OOD).
  - In-context learning: Transformers seem to be good at learning to generalize from data that was not seen during training.
    - Learning to predict the next token when the data is sufficiently complicated requires a general kind of intelligence.
    - Causal inductive bias: The model is biased to predict the next token based on the evidence of the previous tokens.

*Example:* "Based on all the previous evidence, I conclude that the murderer is \_\_\_\_". To do this well, it seems you must be able to reason about the evidence.